



# Scaling Python with Dask

Matthew Rocklin



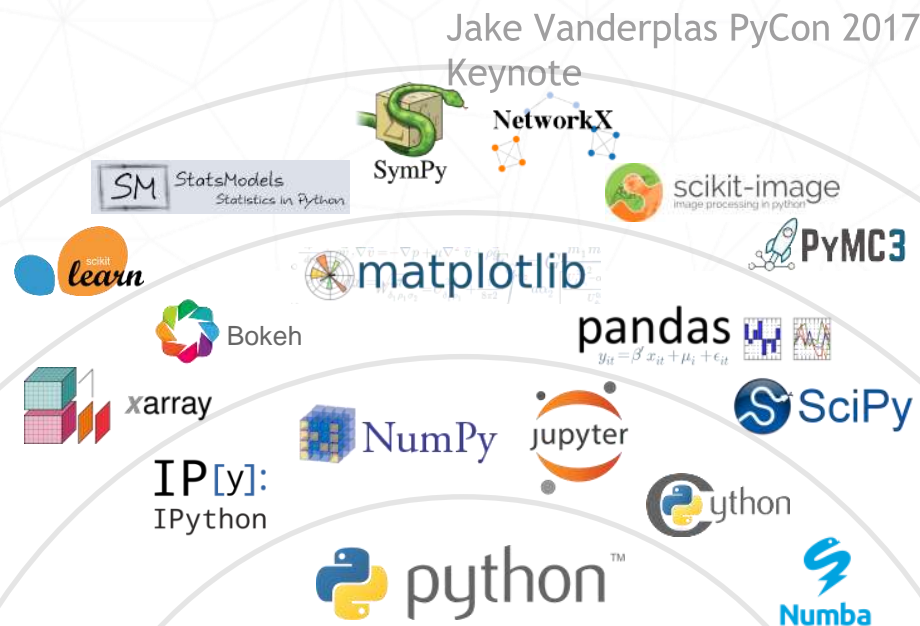
A library for parallel computing in Python

- **Parallelizes** libraries like NumPy, Pandas, and Scikit-Learn
- **Scales** from a personal laptop to large clusters
- **Integrates** easily into existing Python codes
- **Adapts** to custom and sophisticated algorithms

# The Numeric Python Ecosystem

## Capabilities

- **Arrays:** Numpy
- **Dataframes:** Pandas
- **Machine Learning:** Scikit-Learn, ...
- **Interaction:** Jupyter
- **Natural Language:** NLTK, Spacy, Gensim
- **Visualization:** Matplotlib, Bokeh, Altair
- ... thousands of other packages
- ... web, networking, concurrency
- ... mathematics, image processing



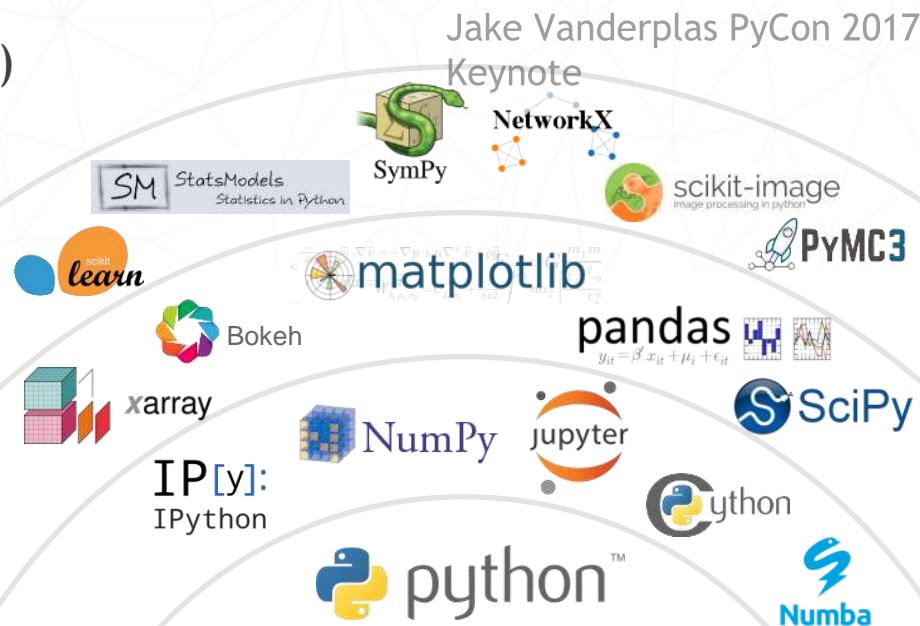
# The Numeric Python Ecosystem

## Strengths

- Easy to use, popular in education
- Fast with both efficient code (C/Fortran) and sophisticated algorithms
- Applied across a wide set of industries

## Weaknesses

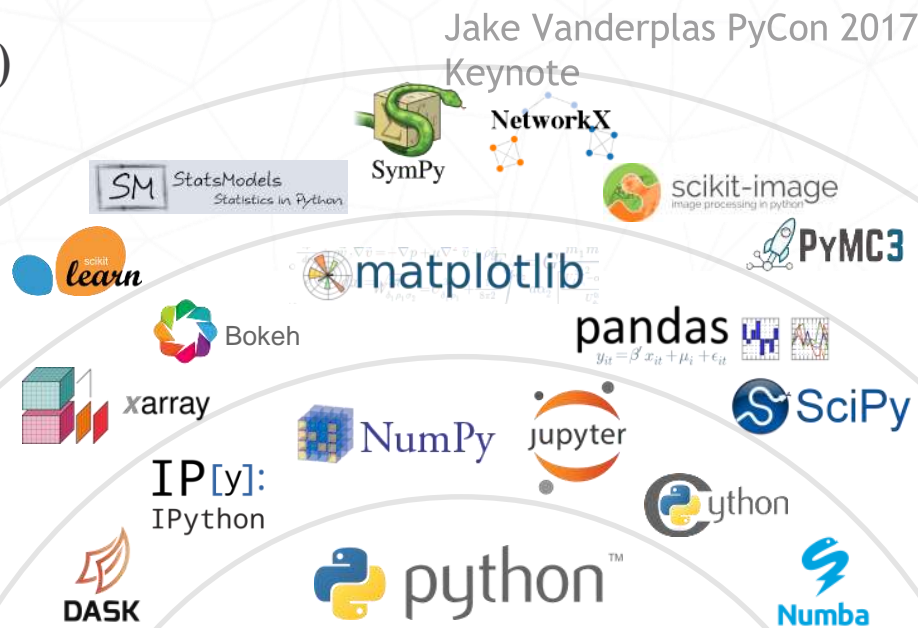
- Designed for in-memory data
- Designed for single-core execution



# The Numeric Python Ecosystem

## Strengths

- **Easy** to use, popular in education
- **Fast** with both efficient code (C/Fortran) and sophisticated algorithms
- **Applied** across a wide set of industries
- Works on **larger-than-memory** data
- Parallelizes across **multi-core** workstations or **distributed** clusters





A library for parallel computing in Python

- **Parallelizes** libraries like NumPy, Pandas, and Scikit-Learn
- **Scales** from a personal laptop to large clusters
- **Integrates** easily into existing Python codes
- **Adapts** to custom and sophisticated algorithms

**High Level:**  
**Parallel Numpy and Pandas**

**Low Level:**  
**Parallelize custom code  
with task scheduling**




# Dask Dataframe

- Scalable implementation of Pandas

Same API (subset)  
Larger datasets

- ```
import pandas as pd
df = pd.read_csv('my-file.csv')
df.groupby('name').value.mean()
```
- Internally coordinates many Pandas dataframes
- Supports reductions, groupbys, joins, timeseries, ...
- Co-evolves with Pandas






# Dask Dataframe

- Scalable implementation of Pandas

Same API (subset)  
Larger datasets

- ```
import dask.dataframe as dd
df = dd.read_csv('s3://bucket/*.csv')
df.groupby('name').value.mean().compute()
```
- Internally coordinates many Pandas dataframes
- Supports reductions, groupbys, joins, timeseries, ...
- Co-evolves with Pandas




# Dask Array

- Scalable implementation of Numpy

Same API (subset)  
Larger datasets

- ```
import dask.array as da
x = da.random.random(...)
y = x.dot(x.T) - x.mean(axis=0)
```

- Internally coordinates many Numpy arrays
- Supports reductions, blockwise, overlapping, linear algebra
- Co-evolves with Numpy



# Demonstration with Dask dataframe

# Parallelizing custom Python code with Dask core

# Parallelize Custom Python Code

- Not all problems are big arrays or dataframes
- Often need fine-grained control
- How would we parallelize this code?

```
for x in A:
    for y in B:
        if x < y:
            z = f(x, y)
        else:
            z = g(x, y)
        results.append(z)
```

- Shouldn't need to rewrite to get parallelism

# Parallelize Custom Python Code

Dask records individual function calls in a task graph

```
def inc(x):  
    return x + 1
```

```
def add(x, y):  
    return x + y
```

```
>>> a = inc(1)  
>>> b = inc(2)  
>>> c = add(a, b)  
>>> c  
5
```


# Parallelize Custom Python Code

Dask records individual function calls in a task graph

```
def inc(x):  
    return x + 1
```

```
def add(x, y):  
    return x + y
```

```
>>> a = inc(1)  
>>> b = inc(2)  
>>> c = add(a, b)  
>>> c  
5
```






# Parallelize Custom Python Code

Dask records individual function calls in a task graph

```
@dask.delayed
def inc(x):
    return x + 1
```

```
@dask.delayed
def add(x, y):
    return x + y
```

```
>>> a = inc(1)
>>> b = inc(2)
>>> c = add(a, b)
>>> c
<Delayed value>
```




# Parallelize Custom Python Code

Dask records individual function calls in a task graph

```
@dask.delayed
def inc(x):
    return x + 1
```

```
@dask.delayed
def add(x, y):
    return x + y
```


```
>>> a = inc(1)
>>> b = inc(2)
>>> c = add(a, b)
>>> c.compute()
5
```



# Parallelize Custom Python Code

Dask records individual function calls in a task graph

```
for x in A:  
    for y in B:  
        if x < y:  
            z = f(x, y)  
        else:  
            z = g(x, y)  
        results.append(z)
```



- Removes barriers to parallelism, used more broadly
- More everyday use of the cluster
- Less rewriting of existing systems


# Demonstration with Task scheduling

**High Level:**  
**Parallel Numpy and Pandas**

**Low Level:**  
**Parallelize custom code  
with task scheduling**

# How does Dask work?


- Dask APIs produce task graphs
- Dask Schedulers execute task graphs



- Dask schedulers can run either ...
  - In a local thread or process pool (lightweight)
  - On a distributed cluster (scalable)

# Dask on a cluster

- Three components
  - Scheduler: centralized metadata
  - Workers: distributed work and storage
  - Clients: like a Jupyter notebook
- Clients interface with users  
Handle APIs like `dask.dataframe`, ...
- Workers do computation  
hold data in distributed memory  
communicate peer-to-peer
- Scheduler coordinates everyone together





# Dask on a cluster

- Manual Setup

```
$ dask-scheduler  
Scheduler at 192.168.0.1:8786
```


```
$ dask-worker 192.168.0.1:8786  
$ dask-worker 192.168.0.1:8786  
$ dask-worker 192.168.0.1:8786
```

```
>>> from dask.distributed import Client  
>>> client = Client('192.168.0.1:8786')
```

- Automatic Setup

- dask-kubernetes
- dask-yarn
- dask-jobqueue (PBS, SLURM, LSF, ...)
- dask-ssh
- ...

- <https://dask.pydata.org/en/latest/setup.html>



# Dask on a laptop

- Dask also works fine without any setup

```
import dask.dataframe as dd
df = dd.read_parquet('my-data.parquet')
result = df.groupby(df.name).balance.sum()

result.compute()    # uses local threads
```

- Uses a local thread pool by default

Easy to get started

Shared memory, very low overhead

# Dask on a Cluster

- Dask also works fine without any setup

```
import dask.dataframe as dd
df = dd.read_parquet('my-data.parquet')
result = df.groupby(df.name).balance.sum()
client = dask.distributed.Client('...')
result.compute() # uses local threads
```

- Now uses a distributed cluster

Easy to scale out when necessary

Cost of switching is low

# Reasons people choose Dask

# 1. Familiar APIs

```
import pandas as pd  
df = pd.read_csv('/path/to/my-file.csv')
```

```
import dask.dataframe as dd  
df = dd.read_csv('hdfs://path/to/my-files-  
*.csv')
```



## 2. Scales Up



- Cloud friendly
- HPC friendly
- Scales to thousands of machines


### 3. Scales Down



- Trivial to install and use on a laptop in a single process
- Smooth transition to multi-core and then distributed operation




## 4. Solves Complex Problems



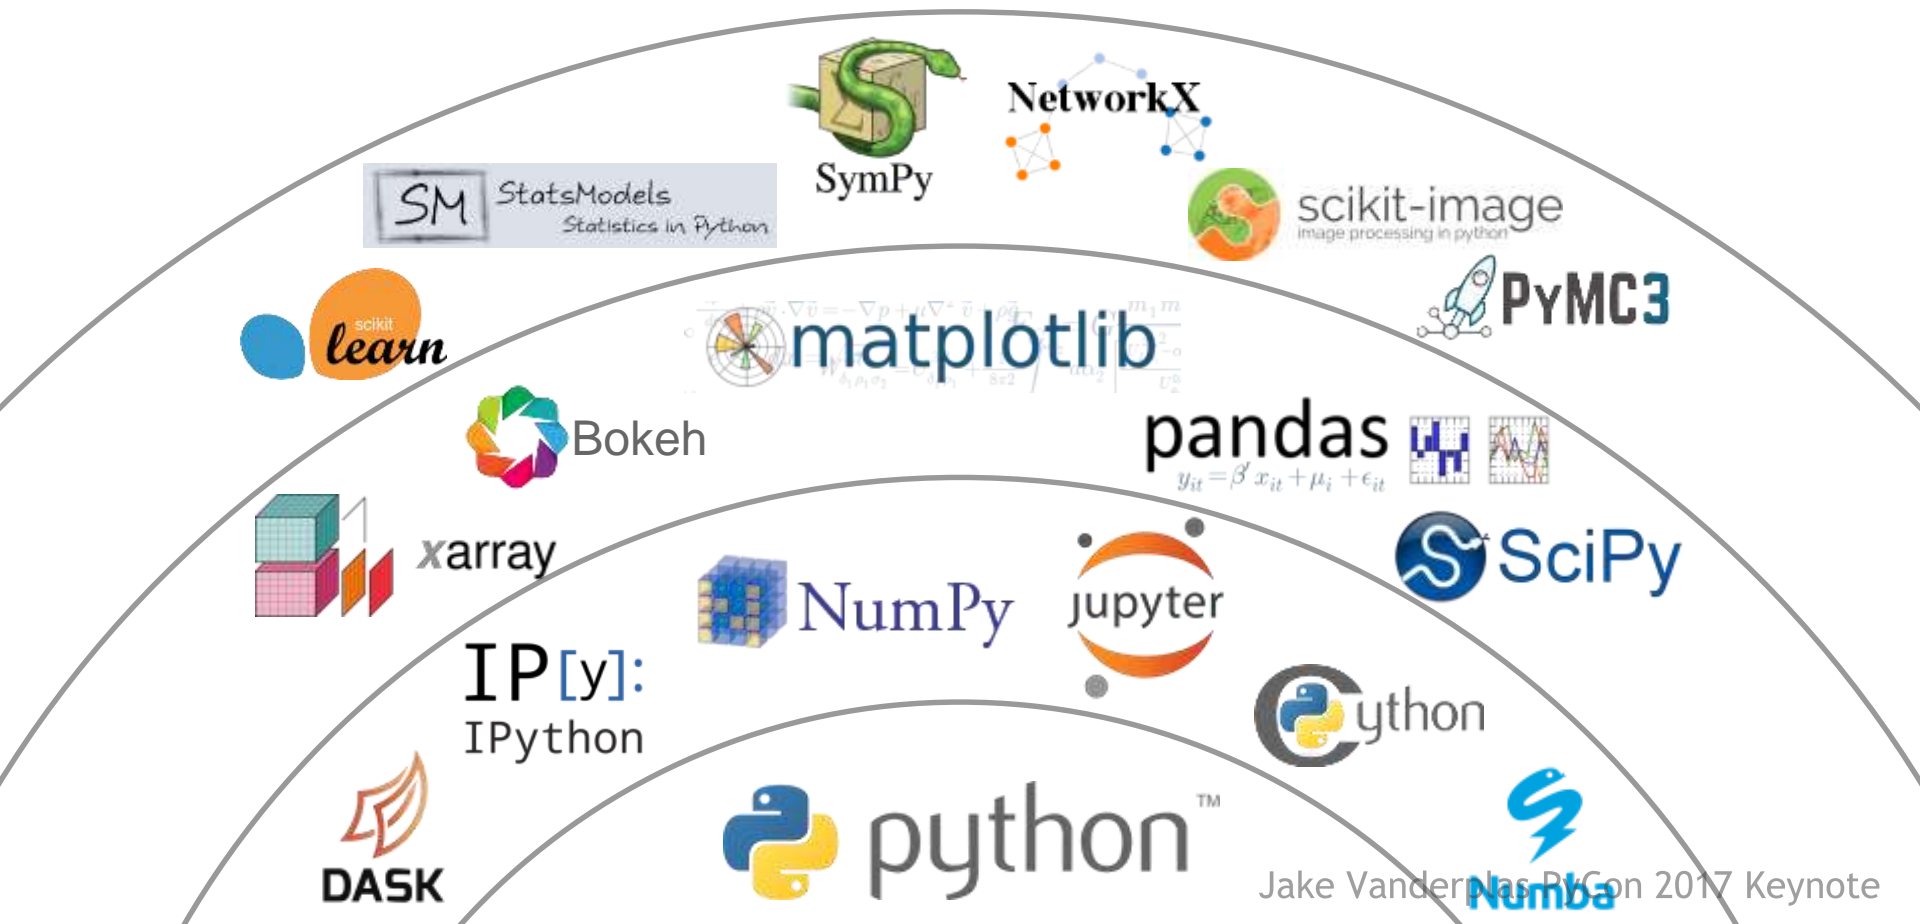
- Many problems are too complex for traditional big data tools
- Dask's scheduler handles these well
- Dask's APIs make it easy to develop bespoke distributed computing systems

## 5. Native Execution



- One of the few non-JVM distributed frameworks
- Plays nicely with native compiled code, GPUs, etc..

## 6. Part of a Broader Ecosystem



# Some Alternatives to Choosing Dask

# Multiprocessing

## Strengths

- Easy to use
- Well understood by many people
- Handles common case problems well

## Weaknesses

- Doesn't handle complex workloads
- Recommend threads when using Numpy/Pandas/Scikit-Learn code

Please also consider `concurrent.futures`

# Apache Spark

## Strengths

- **Implements broad subset of SQL**  
Hooks into BI tooling
- **Integrates well to traditional JVM infrastructure**
- **All-in-one framework**  
You only need to install one thing
- **Well trusted and broadly deployed**

## Weaknesses

- **Doesn't extend well beyond classic tabular computing**
- **JVM-Native code barrier**  
Performance, debugging, usability
- **Reinvents its own ecosystem**  
Doesn't play well with existing systems
- **Often requires wholesale rewrite**



# MPI

## Strengths

- **Very fast**  
still the fastest game in town
- **Very flexible**  
you can implement just about anything
- **Well deployed and supported**  
on high performance computers
- **The only real option for massive parallelism today**

## Weaknesses

- **Hard to use by non-experts**
- **Often overkill for data analysis problems**
- **Not ideal for problems with dynamic load**



# Some Reasons not to Choose Dask

# Dask's limitations

- **Dask is not a SQL database.**  
Does Pandas well, but won't optimize complex queries.  
Consider PostgreSQL, Impala, SparkSQL
- **Dask is not MPI**  
200us task overhead, milliseconds of latency
- **Dask is not a JVM technology**  
Dask targets Python and associated languages (C/C++/...)
- **Dask is not always necessary**  
You may not need parallelism  
Find better algorithms, storage formats, compilers

# Who uses Dask?

# Three main user groups

## 1. People who want big Pandas data frames

Common among early users.  
This is the most common class of questions on Stack Overflow



## 2. Numeric groups with multi-dimensional arrays

Satellite imagery, advanced medical imaging, simulation analysis, signals processing

No other big data system handles this kind of data well

## 3. Advanced groups accelerating their own internal pipelines

Common in finance, operational data acquisition, hardware



# What we didn't talk about

- Scikit-Learn and machine learning
- Array computing
- Real-time applications
- Understanding performance and scalability
- Deploying Dask

# Learn More



A screenshot of a web browser displaying the Dask website. The browser's address bar shows the URL `dask.pydata.org/en/latest/`. The website has a dark theme with orange accents. The navigation bar includes links for "Why Dask?", "Documentation", "Install", "Deploy", and "Tutorial". A dropdown menu under "Documentation" is open, showing links to "Overview", "Arrays", "Dataframes", "Machine Learning", "Custom Applications", and "Real Time". The main content area features the Dask logo (an orange flame-like shape) and the word "DASK" in large white letters. Below the logo, the text "Dask natively scales Python" is displayed, followed by a paragraph: "Dask provides advanced parallelism for analytics, enabling performance at scale for the tools you love". At the bottom of the main content area is a button labeled "Learn More".

← → ↻ ⓘ dask.pydata.org/en/latest/ 🔍 ★ 🔴 🐼 🗨 📄 🗨 📄 📄 📄 📄 📄 📄 📄 📄 📄

Why Dask? Documentation ▾ Install Deploy Tutorial

Overview  
Arrays  
Dataframes  
Machine Learning  
Custom Applications  
Real Time

 **DASK**

Dask natively scales Python

Dask provides advanced parallelism for analytics, enabling performance at scale for the tools you love

[Learn More](#)

# Where to go next



## Download Anaconda

<https://www.anaconda.com/distribution/>



## Test Drive Anaconda Enterprise

[ambassador@anaconda.com](mailto:ambassador@anaconda.com)



## Learn about consulting, training, and support

[ambassador@anaconda.com](mailto:ambassador@anaconda.com)